

Rory Hartong-Redden

October 30, 2024 | Boulder, CO | roryhr@gmail.com | linkedin.com/in/rory-hartong-redden | roryhr.com

Summary

Machine Learning Engineer with over 7 years of experience in data science, data engineering, statistical analysis, and software engineering. Core competencies: Python, ML, and statistical analysis, complemented by 3 years of management experience. Throughout my career I've focused on solving the right problem the right way, understanding my data, writing excellent code, and being a good teammate.

Skills

Programming Languages: Python, SQL, R, Shell DevOps: AWS, Docker, CI/CD, Git, GitHub
Data Processing: Postgres, Spark, Hadoop, S3, Airflow
Python Tools: pandas, scikit-learn, Flask, TensorFlow, Keras, PySpark, Matplotlib, Jupyter

Professional Experience

- **Fast Radius (SyBridge)** Boulder, CO
Technical Manager and Senior Data Scientist Aug 2021–May 2024
 - Led a 3-person data science team, partnering with product and engineering teams to develop and deploy key software products featured in our IPO documents
 - Ensured API uptime of our revenue-critical service, facilitating \$10M in monthly instant quotes
 - Crafted SQL queries to collect data, monitor model accuracy, and answer business questions
 - Trained and deployed a random forest regression model of CNC cycle time for CNC instant quoting, providing the majority of the revenue of our site
 - Provided statistical analysis and visualization analysis for R&D projects with manufacturing IoT time series data
 - **Fast Radius** Chicago, IL
Data Scientist Feb 2020–Aug 2021
 - Tech stack: Python, Flask, Docker, AWS ECS, Datadog, Metabase
 - As the founding data scientist and machine learning engineer I was responsible for building costing and quoting APIs that were backed by regression ML models
 - Predicted shipping costs using mixed integer programming to pack boxes as input for the UPS API
 - In the startup culture I dipped into the Elixir backend, JavaScript frontend, or Terraform infrastructure to fix bugs and remove blockers
 - **runtastic** Linz, Austria
Data Engineer Oct 2018–Sep 2019
 - Tech stack: Python, SQL, Spark, Hadoop, Flume, Oozie, Hive
 - Led the design and deployment of a “People You Might Know” feature to bootstrap a social network and improve the retention rate
 - Built a data exchange prototype with Apache Kafka and a production system with AWS S3
 - Built an ETL pipeline to anonymize customer data to comply with GDPR regulations
 - **Allstate** Menlo Park, CA
Research Analyst Jul 2016–Sep 2018
 - Tech stack: Python, pandas, TensorFlow, Spark, Julia, PostGIS, Keras
 - Trained machine learning models and analyzed telematics and crash data to predict risk
 - Co-authored a research paper titled “Real-time Prediction of Intermediate-Horizon Automotive Collision” with the Stanford Intelligent Systems Lab
-

Education

- **University of California, Santa Barbara**

Santa Barbara, CA

MS Mechanical Engineering

Dec 2014

- Tech stack: MATLAB, SolidWorks, L^AT_EX
- Thesis research: Designed and built an experiment to measure Faraday waves using a novel image processing technique for 3D high-speed mm-resolution measurement over a surface area of 225 cm²

- **University of California, Santa Barbara**

Santa Barbara, CA

BS Physics & BS Mechanical Engineering

June 2010